

## HW3: Model Selection (due 11/22 9:00)

110065508, Cheng-En Lee

### Data Source

```
library(AppliedPredictiveModeling) #install package first!!
library(corrplot) #correlation
library(leaps)
library(glmnet)
library(pls)
library(boot)
```

### Problem 1: Chemical Manufacturing Process Data

#### 1. Data Preprocessing

**(1) Data discription** This data set consists of 177 samples of biological material for which 58 characteristics were measured. Of the 58 characteristics, there were 12 measurements of the biological starting material; 45 measurements of the manufacturing process, and one variable of the resulting final product yield (target variable).

The objective of this problem is to predict the yield given the other variable information.

More data descriptions can be found by typing `help(ChemicalManufacturingProcess)` in r command.

```
data(ChemicalManufacturingProcess)
#View(ChemicalManufacturingProcess) #you may view data in table
#help(ChemicalManufacturingProcess)
#view(dfSummary(ChemicalManufacturingProcess)) #view data summary in html
```

#### (2) Remove NA's

- There are several NA's in the data. We will only use the data with the complete variable information in this homework analysis.
- Rename the dataframe consisting of the complete variable information as **CMP**.
- Rename the variable for a shorten expression

```
dim(ChemicalManufacturingProcess)
```

```
## [1] 176 58
```

```
sum(is.na(ChemicalManufacturingProcess))
```

```
## [1] 106
```

```
which(apply(is.na(ChemicalManufacturingProcess),1,sum)>0) #check which data sample has missing
```

```
## 1 2 3 4 5 6 15 16 17 18 19 20 22 23 24 90 98 134 139 172
## 1 2 3 4 5 6 15 16 17 18 19 20 22 23 24 90 98 134 139 172
## 173 174 175 176
## 173 174 175 176
```

```
CMP <- na.omit(ChemicalManufacturingProcess) #remove missing data
dim(CMP)
```

```
## [1] 152 58
```

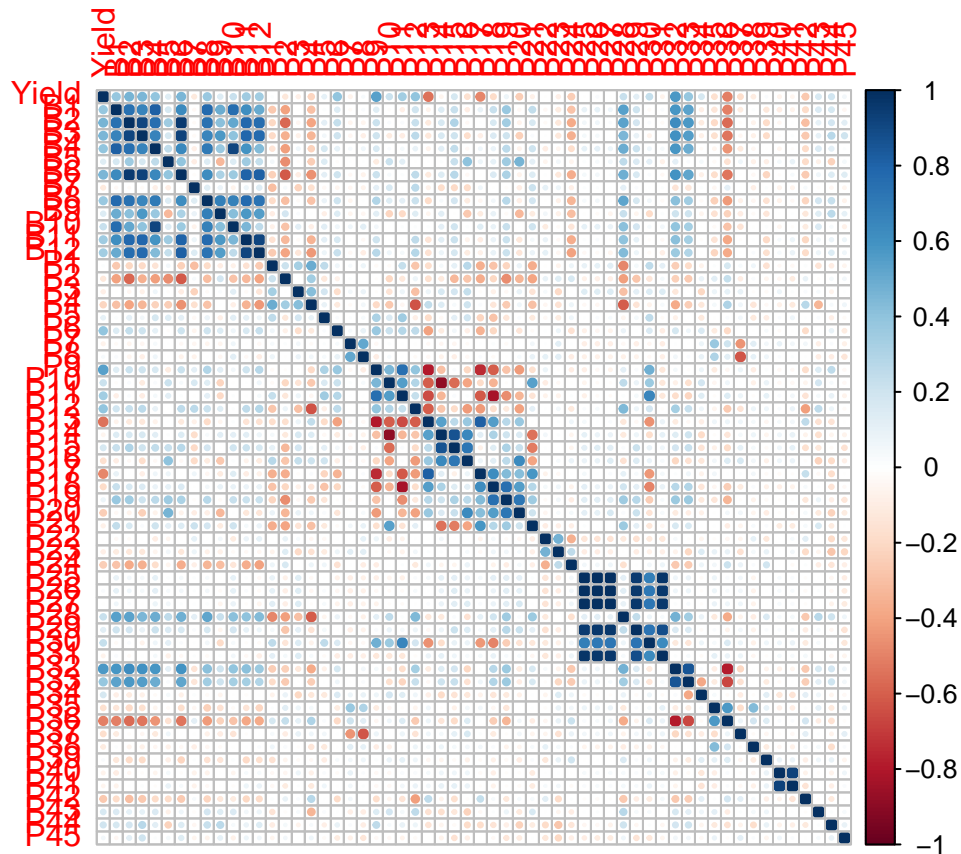
```
sum(is.na(CMP))
```

```
## [1] 0
```

```
#rename variable
B_name = c()
for (i in 1:12){
  B_name[i] = paste("B",i,sep="")
}
P_name = c()
for (i in 1:45){
  P_name[i] = paste("P",i,sep="")
}
names(CMP) <- c("Yield",B_name, P_name)
```

(3) **Correlation plot** Note: pls package also has a function named `corrplot`, which is not we want.

```
corrplot::corrplot(cor(CMP))
```



#### (4) Note

- Fit regression models with various variable selection procedures.
- Go through all variable selection procedures taught in the class, to determine your best model (possibly more than one).
- Identify the important input variables (or any new features you defined) and carefully describe their impacts on the response yield.
- Evaluate your prediction performance via 5-fold cross validation.
- Give a brief analysis summary.

In this homework, I use MSE as the performance metric.

**(5) Bad news: Best Subset does not work** Best subset selection is not feasible for R in this case due to the size of this dataset with the following error:

```
Error in leaps.exhaustive(a, really.big) : Exhaustive search will be S L O W, must specify really.big=T
```

Let's turn to the other algorithms.

**(6) Train test split** The splitting proportion is set to 0.7.

```
set.seed(48763)
train <- sample(c(TRUE, FALSE), nrow(CMP), replace=TRUE, prob=c(0.7,0.3))
test <- !train
```

## 2. Forward Stepwise Selection

(1) **The predict function** In order to conduct CV, we need a predict function since the class `regsubsets` does not have an in-build one.

```
predict.regsubsets <- function(object, newdata, id, ...) {
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id = id)
  xvars <- names(coefi)
  mat[, xvars] %*% coefi
}
```

```
# create a matrix in which we will store the results
k <- 5 # 5-fold CVs
n <- nrow(CMP[train, ])
set.seed(1)
folds <- sample(rep(1:k, length = n))
cv.errors <- matrix(NA, k, 56, dimnames = list(NULL, paste(1:56)))

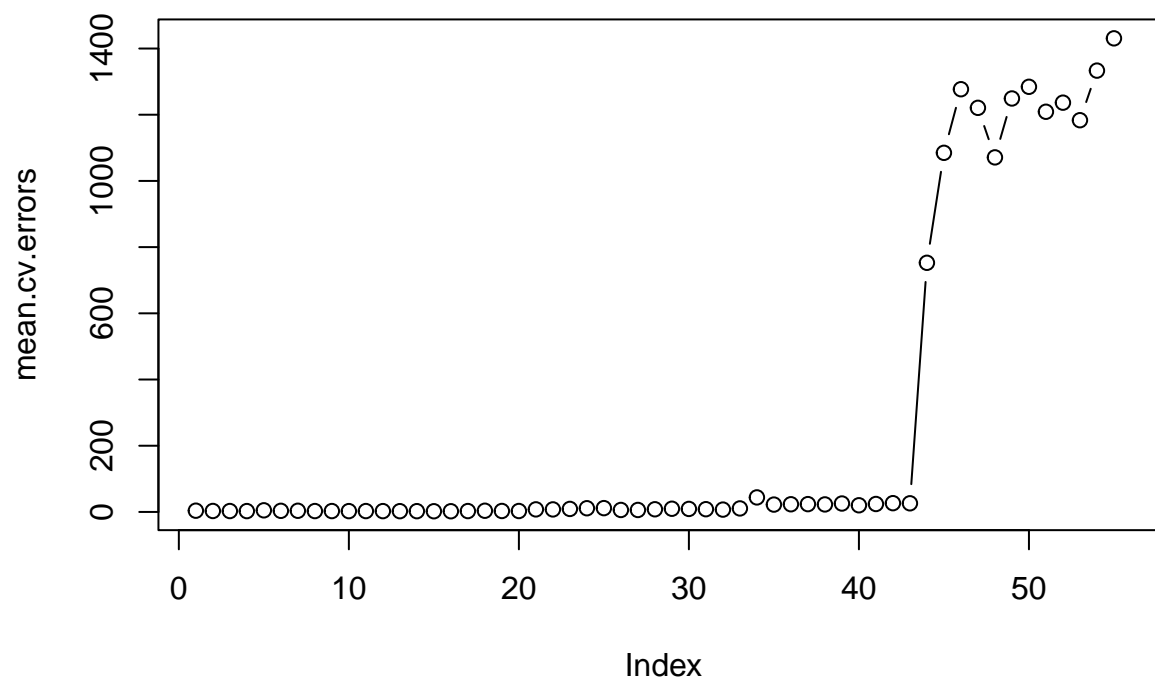
# write a for loop that performs cross-validation
for (j in 1:k) {
  best.fit.fwd <- regsubsets(Yield ~ .,
                            data = CMP[train, ][folds != j, ],
                            nvmax = 56,
                            method = "forward")

  for (i in 1:55) {
    pred <- predict(best.fit.fwd, CMP[train, ][folds == j, ], id = i)
    cv.errors[j, i] <- mean((CMP[train, ]$Yield[folds == j] - pred)^2)
  }
}
```

### (2) 5-fold CV

```
## Reordering variables and trying again:
## Reordering variables and trying again:
## Reordering variables and trying again:
## Reordering variables and trying again:
## Reordering variables and trying again:

# Plot the MSE vs number of features
mean.cv.errors <- apply(cv.errors, 2, mean)
par(mfrow = c(1, 1))
plot(mean.cv.errors, type = "b")
```



Let's find out the best model contains how many features

```
which.min(mean.cv.errors) # How many features does the best model have
```

```
## 16
## 16
```

Since 16 is large enough so that the formula of the model is not suitable to type in the report, I present the coefficients of the model by the following code.

```
reg.best.fwd <- regsubsets(Yield ~ .,
                           data = CMP[train, ],
                           nvmax = 56,
                           method = "forward")
```

## Reordering variables and trying again:

```
coef(reg.best.fwd, 16) # The coefficient
```

```
## (Intercept)      B5      B7      P2      P3
## 203.774714834  0.088596198 -1.803641141  0.013245829 -6.154026209
##      P7      P9      P17      P19      P20
## -0.301997547  0.387444045 -0.405046784  0.012038755 -0.009036993
##      P23      P25      P29      P33      P34
##  0.005510939 -0.002618808  0.646552862  0.292800293  9.455002503
##      P38      P45
## -0.196363005  0.531180811
```

### (3) Model Performance    The training performance

```
mean.cv.errors[16] # training performance
```

```
##          16  
## 2.085848
```

Evaluate the model performance on the testing set

```
test_pred <- predict(reg.best.fwd, CMP[test, ], id = 16)  
MSE <- mean((CMP[test, ]$Yield - test_pred)^2)  
MSE
```

```
## [1] 1.815417
```

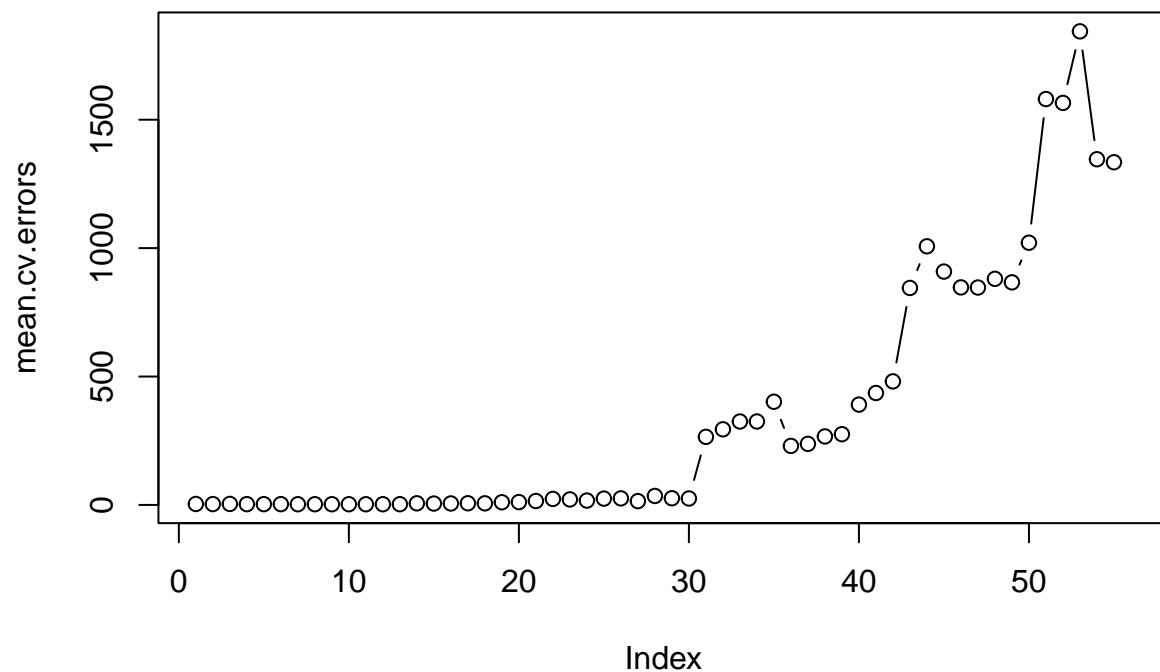
### 3. Backward Stepwise Selection

All the process in this section is parallel to 3.

```
# create a matrix in which we will store the results  
k <- 5 # 5-fold CVs  
n <- nrow(CMP[train, ])  
set.seed(1)  
folds <- sample(rep(1:k, length = n))  
cv.errors <- matrix(NA, k, 56, dimnames = list(NULL, paste(1:56)))  
  
# write a for loop that performs cross-validation  
for (j in 1:k) {  
  best.fit.bwd <- regsubsets(Yield ~ .,  
                             data = CMP[train, ][folds != j, ],  
                             nvmax = 56,  
                             method = "backward")  
  
  for (i in 1:55) {  
    pred <- predict(best.fit.bwd, CMP[train, ][folds == j, ], id = i)  
    cv.errors[j, i] <- mean((CMP[train, ]$Yield[folds == j] - pred)^2)  
  }  
}
```

```
## Reordering variables and trying again:  
## Reordering variables and trying again:  
## Reordering variables and trying again:  
## Reordering variables and trying again:  
## Reordering variables and trying again:
```

```
# Plot the MSE vs number of features  
mean.cv.errors <- apply(cv.errors, 2, mean)  
par(mfrow = c(1, 1))  
plot(mean.cv.errors, type = "b")
```



```
which.min(mean.cv.errors) # How many features does the best model have
```

```
## 11
## 11
```

```
reg.best.bwd <- regsubsets(Yield ~ .,
                           data = CMP[train, ],
                           nvmax = 56,
                           method = "forward")
```

```
## Reordering variables and trying again:
```

```
coef(reg.best.bwd, 11) # The coefficient
```

```
## (Intercept)      B5      P2      P7      P9      P17
## 6.965015596 0.126224215 -0.010264969 -0.357677549 0.418829668 -0.487763889
##      P19      P20      P29      P33      P34      P38
## 0.013464565 -0.008011637 -0.013027450 0.331570014 10.950442953 -0.134254555
```

(3) **Model Performance** The training performance

```
mean.cv.errors[11] # training performance
```

```
##          11  
## 2.785927
```

Evaluate the model performance on the testing set

```
test_pred <- predict(reg.best.bwd, CMP[test, ], id = 11)  
MSE <- mean((CMP[test, ]$Yield - test_pred)^2)  
MSE
```

```
## [1] 1.826801
```

Note: It looks like there are some bugs in `regsubsets` since I cannot choose every features into the model. To be precisely, the argument `nvmax` can only set to 56. However, this does not affect the results since it is really unlikely that the best model contains “every” features.

#### 4. Ridge regression

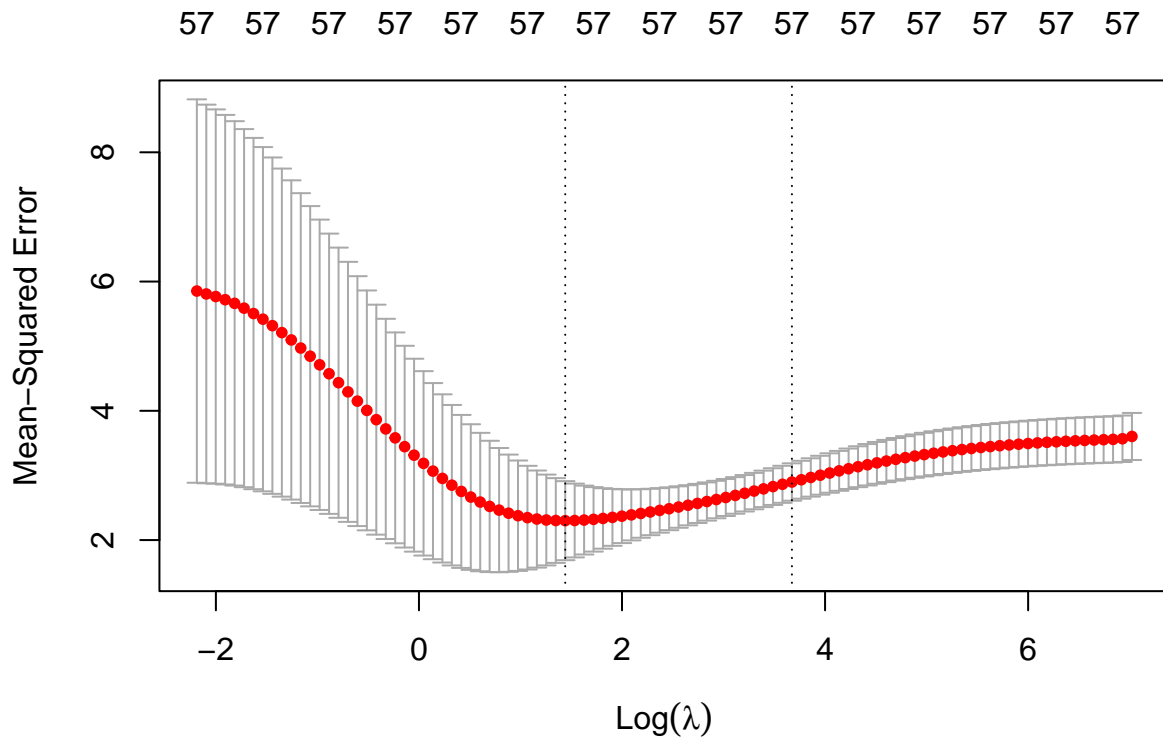
```
x.train <- model.matrix(Yield ~ ., CMP[train, ])[, -1]  
y.train <- CMP[train, ]$Yield  
x.test <- model.matrix(Yield ~ ., CMP[test, ])[, -1]  
y.test <- CMP[test, ]$Yield
```

(1) Divide the target variable and features

(2) In-built CV In the `glmnet` function, if the argument `alpha=0` then a ridge regression model is fit, and if `alpha=1` then a lasso model is fit.

```
grid <- 10^seq(10, -2, length = 100) # use grid search to find lambda  
set.seed(48763)  
cv.out <- cv.glmnet(x.train, y.train, alpha = 0) # Ridge regression  
plot(cv.out)
```





By the following code, we know that the best  $\lambda \approx 4.635$ .

```
bestlam_ridge <- cv.out$lambda.min # best lambda
bestlam_ridge
```

```
## [1] 4.223584
```

**(3) Model Performance** Evaluate the performance of Ridge regression.

```
# retrain the model with the best lambda
ridge.mod <- glmnet(x.train, y.train, alpha = 0, lambda = grid)

# training performance
ridge.pred <- predict(ridge.mod, s = bestlam_ridge, newx = x.train)
MSE_train <- mean((ridge.pred - y.train)^2)

# testing performance
ridge.pred <- predict(ridge.mod, s = bestlam_ridge, newx = x.test)
MSE_test <- mean((ridge.pred - y.test)^2)

# results
MSE_train
```

```
## [1] 1.370079
```

```
MSE_test
```

```
## [1] 1.557925
```

(4) **Important Features** Since Ridge regression does not really select variables, I show the coefficient of the first 20 features.

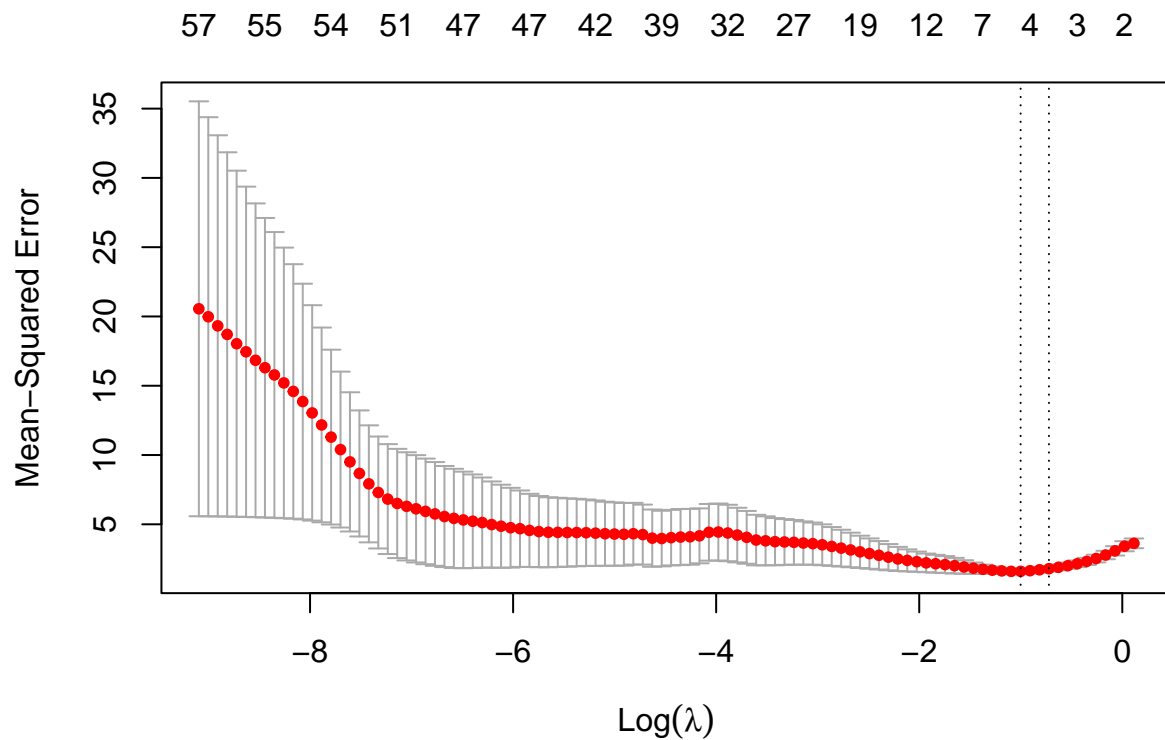
```
ridge_coeff <- predict(ridge.mod,
                        s = bestlam_ridge,
                        exact = T,
                        type = "coefficients",
                        x = x.train,
                        y = y.train)
ridge_coeff[1:20, ]
```

|    |               |               |               |              |               |
|----|---------------|---------------|---------------|--------------|---------------|
| ## | (Intercept)   | B1            | B2            | B3           | B4            |
| ## | 71.8476344121 | 0.0643033660  | 0.0155283234  | 0.0131913681 | 0.0254978603  |
| ## | B5            | B6            | B7            | B8           | B9            |
| ## | 0.0119775822  | 0.0132245521  | -0.4925201600 | 0.0566020940 | -0.0060573688 |
| ## | B10           | B11           | B12           | P1           | P2            |
| ## | 0.0158912180  | 0.0090817035  | 0.0368106962  | 0.0001256437 | 0.0001350158  |
| ## | P3            | P4            | P5            | P6           | P7            |
| ## | -1.5131671132 | -0.0017485545 | 0.0003365234  | 0.0595367295 | -0.0841868700 |

## 5. Lasso

The process in this section is parallel to 5.

```
set.seed(1)
cv.out <- cv.glmnet(x.train, y.train, alpha = 1) # LASSO
plot(cv.out)
```



By the following code, we know that the best  $\lambda \approx 0.3673$ .

```
bestlam_lasso <- cv.out$lambda.min # best lambda
bestlam_lasso
```

```
## [1] 0.3673458
```

```
# retrain the model with the best lambda
lasso.mod <- glmnet(x.train, y.train, alpha = 1, lambda = grid)

# training performance
lasso.pred <- predict(lasso.mod, s = bestlam_lasso, newx = x.train)
MSE_train <- mean((lasso.pred - y.train)^2)

# testing performance
lasso.pred <- predict(lasso.mod, s = bestlam_lasso, newx = x.test)
MSE_test <- mean((lasso.pred - y.test)^2)

# results
MSE_train
```

```
## [1] 1.45914
```

```
MSE_test
```

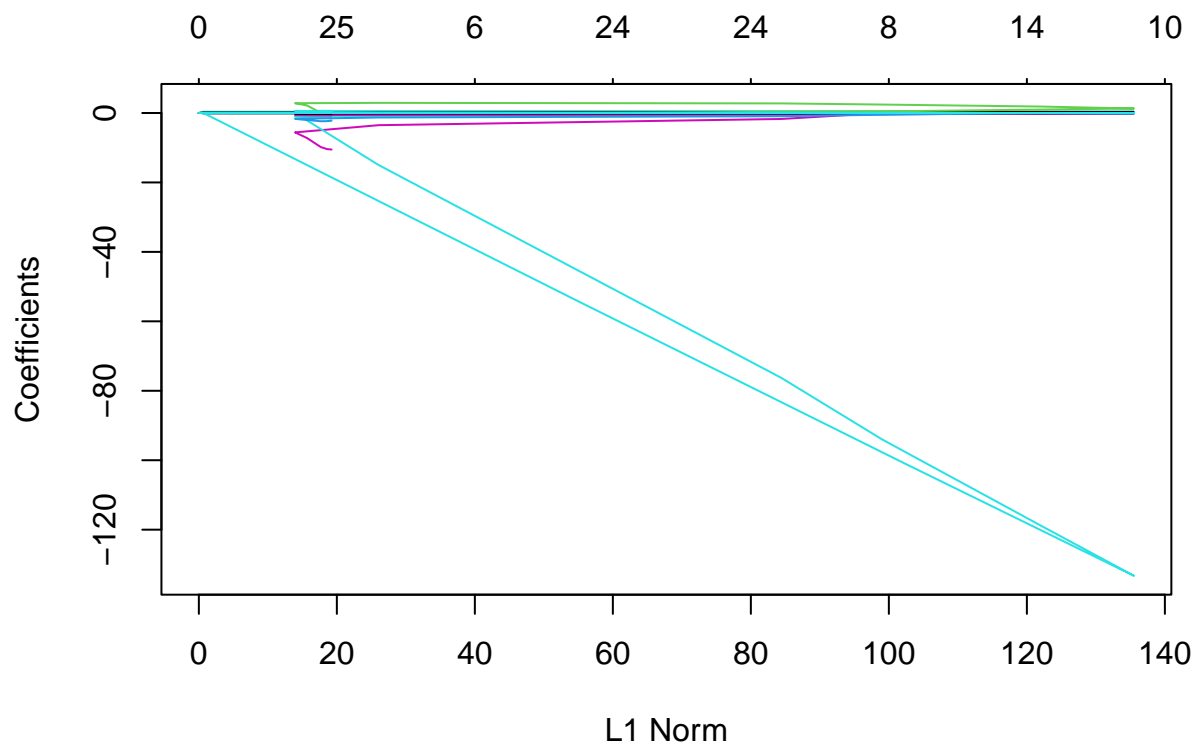
```
## [1] 1.695445
```

(4) **Important Features** It turns out that LASSO selects 5 features only.

```
lasso_coeff <- predict(lasso.mod,  
                        s = bestlam_lasso,  
                        exact = T,  
                        type = "coefficients",  
                        x = x.train,  
                        y = y.train)  
lasso_coeff[1:57,][which(lasso_coeff!=0)]
```

```
## (Intercept)          P6          P9          P13          P17          P32  
## 4.846384160 0.004737659 0.367678143 -0.097801125 -0.049232780 0.142601385
```

```
plot(lasso.mod)
```



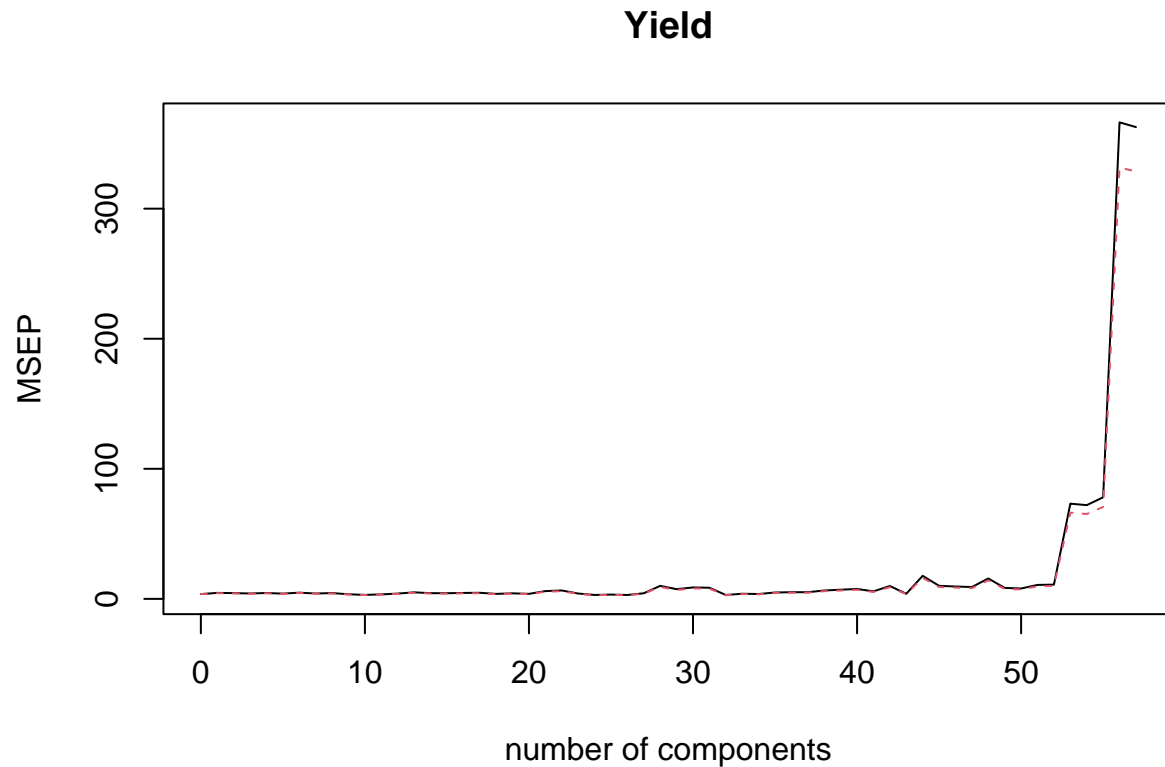
## 6. PCA

Perform PCR on the training data.

```

set.seed(48763)
pcr.fit <- pcr(Yield ~ .,
               data = CMP,
               subset = train,
               scale = TRUE,
               validation = "CV")
validationplot(pcr.fit, val.type = "MSEP")

```



```

MSEP_object <- MSEP(pcr.fit)
which.min(MSEP_object$val[1,1, ])

```

```

## 32 comps
##      33

```

Hence PCR selects 33 features. Also, model performance can be evaluated

```

# training performance
pcr.pred.train <- predict(pcr.fit, x.train, ncomp = 33)
MSE_pcr_train <- mean((pcr.pred.train - y.train)^2)

# testing performance
pcr.pred.test <- predict(pcr.fit, x.test, ncomp = 33)
MSE_pcr_test <- mean((pcr.pred.test - y.test)^2)

```

```
# results
MSE_pcr_train
```

```
## [1] 0.7671907
```

```
MSE_pcr_test
```

```
## [1] 1.689532
```

## 7. Summary.

**(1) Model performance** By the above discussion, the best model is Ridge regression, considering the testing MSE only. Both FSS and BSS are not performs as well as Ridge. However, considering the simplicity of the model, perhaps LASSO is more feasible since it only takes 5 features into consideration.

```
knitr::include_graphics("1.png")
```

| Model | Num. of feat. | Training MSE | Testing MSE |
|-------|---------------|--------------|-------------|
| FSS   | 16            | 2.0858       | 1.8154      |
| BSS   | 11            | 2.7859       | 1.8268      |
| Ridge | X             | 1.3701       | 1.5579      |
| LASSO | 5             | 1.4591       | 1.6954      |
| PCR   | 33            | 0.7672       | 1.6895      |

Though PCR performs slightly better than LASSO, the gap between training MSE and testing MSE suggests that there is an overfitting. The same problems occurs when it comes to BSS.

**(2) Important variables** The below table shows the important variables select by FSS, BSS, and LASSO. Only P9 and P17 is regarded important by all of the 3 algorithms. The coefficient in LASSO suggests that P9 has a positive effect on Yield, while P17 has a negative effect.

```
knitr::include_graphics("2.png")
```

| FSS | BSS | LASSO |
|-----|-----|-------|
| B5  | B5  |       |
| B7  |     |       |
| P2  | P2  |       |
| P3  |     |       |
|     |     | P6    |
| P7  | P7  |       |
| P9  | P9  | P9    |
|     |     | P13   |
| P17 | P17 | P17   |
| P19 | P19 |       |
| P20 | P20 |       |
| P23 |     |       |
| P25 |     |       |
| P29 | P29 |       |
|     |     | P32   |
| P33 | P33 |       |
| P34 | P34 |       |
| P38 | P38 |       |
| P45 |     |       |

---

## Problem 2: Data with Outliers

For iid data  $X_i, i = 1, 2, \dots, n$  from some distribution, the sample standard deviation  $\hat{\sigma}$  is often used to estimate the population deviation  $\sigma$ :

$$\hat{\sigma} = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2}.$$

But sometimes, the observed data are contaminated with extreme values or outliers, due to recording errors or other reasons. In order to obtain a robust estimate for  $\sigma$ ,  $\tilde{\sigma}$  is considered:

$$\tilde{\sigma} = 1.4826 \cdot \text{med}_{1 \leq i \leq n} \{|X_i - X_{\text{med}}|\},$$

where  $X_{\text{med}}$  denotes the sample median of  $\{X_i : 1 \leq i \leq n\}$ .

Apply the nonparametric bootstrap method to evaluate the sampling distributions of  $\tilde{\sigma}$  and  $\hat{\sigma}$ , and their estimation precision  $\text{var}(\tilde{\sigma})$  and  $\text{var}(\hat{\sigma})$ .

The data for problem 2: **hw3\_problem2.csv**

```
data <- read.table("hw3_problem2.csv", sep = ",", header=TRUE)[,2]
```

## 1. The function to evaluate the deviations

```
sigma_hat <- function(dat, index) {  
  X <- dat[index]  
  sqrt((1/(length(dat)-1)) * (sum( (X-mean(X))^2 )))  
}  
  
sigma_tilde <- function(dat, index) {  
  X <- dat[index]  
  1.4826*median(abs(X-median(X)))  
}
```

## 2. Bootstarp

```
boot(data, sigma_hat, R = 1000)
```

```
##  
## ORDINARY NONPARAMETRIC BOOTSTRAP  
##  
##  
## Call:  
## boot(data = data, statistic = sigma_hat, R = 1000)  
##  
##  
## Bootstrap Statistics :  
##      original      bias    std. error  
## t1* 3.745385 -0.0273867   0.3565152
```

```
boot(data, sigma_tilde, R = 1000)
```

```
##  
## ORDINARY NONPARAMETRIC BOOTSTRAP  
##  
##  
## Call:  
## boot(data = data, statistic = sigma_tilde, R = 1000)  
##  
##  
## Bootstrap Statistics :  
##      original      bias    std. error  
## t1* 2.585234 0.05279911   0.2867246
```

Hence we have

$$\hat{\sigma} \approx 3.7454, \text{Var}(\hat{\sigma}) \approx 0.3698$$

and

$$\tilde{\sigma} \approx 2.5852, \text{Var}(\tilde{\sigma}) \approx 0.2987$$